




Krrish Khanna | Technical Operator & Founder's Office Portfolio

Technical Operator | AI-Native Builder | Founder's Office Aspirant

"I am obsessed with leverage, systems, and building things that scale."

📍 Chennai, India 🇮🇳 · 🎓 SRM IST × Waseda University 🇯🇵

🔧 **Co-founder of Odin** - co-building a live subtitles & real-time translation app (100+ languages) with a friend, shipped with a live Free / Pro / Enterprise model.

⚡ **"I sit at the intersection of technical execution, data analytics, and business strategy. I am an entrepreneurial generalist with a CS background who uses AI-native workflows to build, fail, learn, and iterate at 10x velocity. I don't just write code - I solve high-ambiguity business problems."**



[!\[\]\(529949c2c3dadbaa4e538e8c643454bc_img.jpg\) LinkedIn](#) | [!\[\]\(9d83b67c094360bb3c4e3b68ca3d779f_img.jpg\) GitHub](#) | [!\[\]\(7dbe5b492efc9d2ec2df517769c7fbf7_img.jpg\) krrishkhanna08@gmail.com](mailto:krrishkhanna08@gmail.com) | [!\[\]\(3ace734828313c58d082f679b9285da0_img.jpg\) Resume](#)

Selected Wins



- **Co-founded & shipped Odin** - a live subtitles & translation app (Windows + Android, 100+ languages) co-built with a friend, with Free / Pro / Enterprise pricing and a live go-to-market.
- **91% validation accuracy** across 120+ test scenarios - VeriLoad shipped solo, end-to-end.
- **65% faster query-to-insight** - AI knowledge-retrieval workflows at Git Software Technologies.
- **1.5x-2x inference speedup & 40% lower memory** on compressed LLMs at Samsung R&D.
- **4-5 lakh records** ingested through a validation pipeline adopted by 10-20 stakeholders at KPMG.
- **28% precision gain** across 15,000+ multilingual claims on Vachan.



Current Focus



Currently exploring opportunities in:

- Founder's Office
- Technical Operator
- Product & Strategy
- AI-Native Startups

I enjoy operating at the intersection of technology, business, and execution.

Long-term, I aspire to build companies. **Short-term**, I want to learn how exceptional companies are built from the inside.



What I'm Looking For



I'm obsessed less with *startups* and more with **leverage, agency, ownership, and systems** - startups are simply the best vehicle for all four. I do my best work when given a difficult problem, high ownership, limited instructions, and the freedom to figure things out.



Why Founder's Office?



I've never been particularly interested in staying within one function. Throughout university, I explored AI, analytics, consulting, product development, systems design, and business strategy. What fascinated me wasn't the individual disciplines - it was how they connected.

I naturally gravitate toward questions like:

- What's the bottleneck?
- What's the highest-leverage action?
- What system is producing this outcome?
- How can we scale this?

The Founder's Office role appeals to me because it sits at the intersection of all these questions. I enjoy ambiguity, ownership, and operating across functions.

I don't want to be handed a playbook. I want to help build one.



How I'm Wired

I've always been more interested in understanding systems than memorizing answers. Whether it's technology, business, markets, products, organizations, or people, my instinct is usually the same:

How does this work? What's the underlying system producing this outcome?

I naturally gravitate toward connecting ideas across disciplines rather than staying within a single domain. Over time, this has led me to explore AI, analytics, product

development, strategy, finance, and operations - not because I wanted to collect skills, but because I was trying to understand how the pieces fit together.

I tend to think in terms of **leverage**. When faced with a problem, my first instinct isn't "How do we work harder?" - it's usually:

- What's the bottleneck?
- What's the highest-leverage action?
- Can this be automated?
- Is there a better system?

I enjoy operating in ambiguity. I don't need every step defined before I begin - I often find the most interesting work exists where the path isn't obvious yet.

Ownership energizes me. I do my best work when I'm trusted with a problem, given room to think, and allowed to drive it from idea to execution.

I'm also relentlessly curious. I dive into topics far outside my immediate field - from startup strategy and consumer behavior to financial markets, AI systems, and product design. The common thread is a desire to understand how complex systems create outcomes.

At my core, I'm less interested in being the best at a single function and more interested in becoming exceptionally good at **solving problems that sit between functions**. That's the work I find most meaningful.

How I Solve Problems

My default approach when faced with ambiguity:

1. **Understand the desired outcome** - what does "done" actually look like?
2. **Identify the bottleneck** preventing that outcome.
3. **Gather only the information** required to make the decision - no more.
4. **Prioritize by leverage, not effort.**
5. **Execute quickly.**
6. **Measure results.**
7. **Iterate.**

The framework in action:



KPMG - *Outcome:* enterprise adoption of a classification pipeline under regulatory ambiguity. *Bottleneck:* no agreed success metrics across 3 client teams. *Decision:* I defined the latency / accuracy / adoption metrics myself and drove the go/no-go calls instead of waiting for sign-off. *Result:* adoption across 10-20 stakeholders and 4-5 lakh records ingested.



VeriLoad - *Outcome:* a working dispatch-validation product, solo. *Bottleneck:* unbounded scope with no spec. *Decision:* I built and validated the core scoring engine first and deferred all UI polish until accuracy hit 91%. *Result:* shipped a live stakeholder-facing dashboard end-to-end.



Samsung - *Outcome:* faster, cheaper LLM inference. *Bottleneck:* GPU memory and latency on compressed models. *Decision:* I prioritized isolating the network-contraction bottleneck over broad, unfocused experimentation. *Result:* a 1.5x-2x inference speedup with ~40% lower memory.



Decisions, Lessons & Beliefs

Lessons I've learned the hard way:



Building fast > building perfect. I used to over-optimize early versions. Most of the useful learning only showed up *after* shipping.



Users rarely care about your architecture. I once chased technical elegance; I've learned users only care whether the problem actually gets solved.



Data doesn't replace judgment. Analytics informs a decision - it doesn't make it. Someone still has to own the call.

A few things I believe:

- Most students optimize *credentials* instead of *capability*.
- AI will increase the value of judgment, not decrease it.
- Speed is chronically underrated inside large organizations.
- Curiosity compounds faster than talent.



Section 1: Shipped Products & AI-Native Prototyping

This section highlights my ability to act as a solo builder, taking vague concepts and turning them into functional, user-facing software without predefined specs.



Flagship - Odin: the product I co-founded and am building with a friend - my most complete 0-to-1 work. A system-wide live-subtitles and real-time-translation overlay (Windows + Android, 100+ languages) - shipped *with* a real go-to-market: Free / Pro / usage-based Enterprise tiers, hosted cloud speech, and a Discord + Instagram community motion. Full breakdown in the Shipped Products gallery below.



VeriLoad - Solo-Shipped Dispatch Platform

- **The Context:** Shipped a working B2B product solo end-to-end using advanced AI coding tools.

- **The Problem:** Manual dispatch reconciliation environments are highly ambiguous, error-prone, and lack real-time validation frameworks.
- **The Execution:** Led independent product discovery, requirements definition, and prototype development. Engineered a statistical validation engine with **91% accuracy across 120+ test scenarios** and launched a live stakeholder-facing dashboard.



Operator Value: Proves high agency, rapid prototyping capabilities, and the execution grit required to build systems out of pure chaos.



Vachan - Multilingual Fact-Checking Platform

- **The Context:** Multi-source NLP and data pipeline handling over **15,000+ information claims**.
- **The Execution:** Designed and deployed a multi-source processing pipeline. Iterated aggressively with LLM and AI tools to improve model precision by **28% over the baseline**, and built a live dashboard for real-time accuracy and throughput reporting.



Operator Value: Proves heavy data literacy and the ability to translate complex engineering backend pipelines into clean, transparent business metrics.



Section 2: Corporate & Research Execution

This section highlights my experience operating in high-stakes environments, managing senior stakeholders, and optimizing large-scale operational infrastructure.



KPMG India - Technical Trainee Intern (Aerospace & Defence)

- **The Business Problem:** A high-value corporate client required an end-to-end classification system wrapped in deep regulatory ambiguity.

- **The Execution:** Operated completely unsupervised. Ran parallel discovery tracks with **3 distinct client teams** to scope requirements, defined success metrics (latency, accuracy, adoption), and independently drove critical go/no-go calls with senior leadership.
- **The Impact:** Successfully drove operational adoption of the resulting pipeline across **10-20 enterprise stakeholders** for the seamless ingestion of **4-5 lakh data records**.
- **What I Learned:** Stakeholder management, navigating regulatory ambiguity, and validating enterprise-scale data.

Samsung Research Institute - R&D Operator (LLM Infrastructure Optimization)

- **The Business Problem:** Deploying state-of-the-art Large Language Models (like Llama-2-7B) introduces massive infrastructure friction, heavy GPU memory demands, and slow inference latency.
- **The Execution:** Configured and optimized tensorized LLM environments using Activation-Aware SVD on high-performance computing clusters (NVIDIA A100 GPUs). Engineered an efficient, iterative calibration process to optimize layer-specific decomposition, and explored custom model extensions for CPU simulations to minimize cloud compute dependencies.
- **The Impact:** Identified core network contraction bottlenecks, directly contributing to a major **>1.5x to 2x inference performance speedup** for compressed LLMs.
- **What I Learned:** Performance optimization, operational data analysis, and reasoning about engineering tradeoffs.



Operator Value: Proves I understand the unit economics of AI. I focus directly on infrastructure cost and execution speed-the two metrics early-stage AI founders care about most.

Git Software Technologies - AI Engineering & Operations Intern

- **The Business Problem:** Distributed, international teams face severe operational drag due to fragmented internal knowledge bases and disconnected enterprise data sources.
- **The Execution:** Acted as an autonomous operator to build modular API connectors bridging multiple enterprise data sources, adapting existing code bases without waiting for a handed-down spec. Scoped and shipped AI-powered knowledge retrieval workflows utilizing Cursor and Claude, iterating directly from stakeholder feedback.
- **The Impact:** Successfully reduced query-to-insight time by **65%**. Defined and benchmarked performance metrics independently, presenting the final business impact directly to distributed international stakeholders.
- **What I Learned:** AI-native development, rapid prototyping, and operating across distributed teams.



Operator Value: Proves an AI-native execution premium. It shows I can independently build the automated workflows needed to keep a global team running fast and lean.

Section 3: Global Perspective & Achievements

Rigorous academic standards and a global outlook developed through international immersion.



Having lived, studied, and travelled across India, Japan, Singapore, and Australia, I started paying attention to *what each place taught me* - not just that I'd been there:

- **Japan taught me about systems** - how precise, well-designed systems remove the need for heroics.
- **Australia taught me about autonomy** - lean, product-led teams doing more with less.
- **India taught me about scale** - building for extreme price-sensitivity, volume, and infrastructure constraints.
- **Singapore taught me about trust** - how standards and infrastructure create trust at scale.

The travel only matters because of the insight it left behind.

- **Waseda University (Tokyo, Japan):** Selected for a competitive Semester Exchange within the School of Fundamental Science and Engineering.
- **Leadership Initiatives:** Selected for the McKinsey Forward Program and the Harvard Faculty-led Aspire Leaders Program, honing structured business problem-solving and strategic leadership.



Section 4: Global Product Strategy & Case Studies

This section demonstrates how I analyze market patterns, evaluate user retention, and build frameworks for hyper-growth across different global tech hubs.



The Method: End-to-end product teardown analyses across consumer technology and AI markets - evaluating user journeys, retention loops, and monetization strategies, then translating them into feature-prioritization frameworks that weigh user value against implementation complexity.

Indian Unicorns: High-Frequency Operational & Transactional Scale

Swiggy Instamart vs. Zomato Blinkit - The Quick-Commerce Unit Economics War

- **The Core Conflict:** Winning the hyper-local 10-minute delivery market through dark store optimization, density, and average order value (AOV) expansion.
 - **The User Retention Loop:** Both ecosystems utilize high-value subscription lock-ins (Zomato Gold vs. Swiggy One) to lower transaction friction. By eliminating delivery fees, they turn low-frequency discretionary buyers into high-frequency habitual users, heavily increasing Customer Lifetime Value (LTV).
 - **Product Strategy Analysis:**
 - **Dark Store Geometry & SKU Control:** Blinkit's superior supply chain execution relies on strict real-time prediction of neighborhood demand, optimizing local SKU selection to maximize inventory turnover.
 - **The Upsell Funnel:** Analyzing the checkout journey reveals aggressive micro-conversions - adding low-cost "impulse buy" items (snacks, seasonal products) at checkout to push the basket size above profitability thresholds.
-  **Operator Takeaway:** Success in this space isn't driven by code aesthetic; it is won by narrowing algorithmic routing latency and maximizing the margin per square foot of real estate.



If I were running Blinkit:

- **3 priorities:** sharpen per-dark-store demand forecasting; lift AOV through checkout upsells; defend contribution margin per order.
- **3 risks:** cash burn from over-expansion; rider and throughput cost inflation; subscription fatigue eroding loyalty.
- **3 opportunities:** high-margin private-label SKUs; retail-media and ad revenue; hyperlocal SKU personalization by neighborhood.



ixigo - High-Utility Monetization for Next-Billion Users

- **The Core Conflict:** Dominating travel tech in Tier-2 and Tier-3 Indian markets where internet bandwidth is volatile and price sensitivity is extreme.
- **The User Retention Loop:** ixigo positioned its app as a daily utility (free offline train running status and PNR tracking) rather than a seasonal booking engine. This low-cost utility keeps monthly active users (MAU) high without requiring continuous ad-spend.
- **Product Strategy Analysis:**
 - **Frictionless UX:** Built an ultra-lightweight app structure that caches critical transit data locally on the user's phone.
 - **Micro-Insurance Monetization:** Since users are price-conscious, ixigo drives profitability through high-margin ancillary features, such as one-tap "Instant Free Cancellation" insurance options embedded directly into the payment flow.



Operator Takeaway: To capture emerging consumer markets, build a high-frequency free utility first, then monetize through seamless, low-friction financial add-ons.



If I were Founder's Office at ixigo:

- **3 priorities:** deepen the free-utility habit loop; grow high-margin ancillaries (insurance, cancellation); keep the app ultra-light for low-bandwidth users.
- **3 risks:** razor-thin margins on core ticketing; OTA price wars; over-reliance on the rail vertical.
- **3 opportunities:** embedded fintech and micro-insurance; regional-language expansion; a Tier-2/3 brand-trust moat.



Australian Tech: Mastering Product-Led Growth (PLG)



Canva - Scaling Enterprise Value Without Losing Simplicity

- **The Core Conflict:** Moving upstream from casual, freemium creators into highly lucrative enterprise B2B workflows while maintaining a dead-simple interface.
- **The User Retention Loop:** Viral collaboration loops. When an individual designer shares a project link with a team member, the recipient must create a Canva profile to collaborate. This turns the product itself into the primary user acquisition engine (lowering CAC).
- **Product Strategy Analysis:**
 - **AI-Native Extension (Magic Studio):** Canva successfully integrated advanced generative AI editing assets into its interface without confusing its non-technical user base. They achieved this by masking complex prompt engineering behind simple, familiar drag-and-drop actions.
 - **Workspace Monopolization:** Creating shared brand kits, centralized font books, and team asset folders makes it operationally painful for an organization to migrate to a competitor.



Operator Takeaway: True Product-Led Growth requires building native virality directly into the core user loop so that user expansion happens organically without a bloated sales force.



If I were PM at Canva:

- **3 priorities:** convert free creators into paid Teams; embed Magic Studio AI without UX bloat; deepen brand-kit lock-in.
- **3 risks:** AI design tools commoditizing the core; enterprise security and compliance demands; feature sprawl diluting simplicity.
- **3 opportunities:** vertical template packs; collaboration virality into new orgs; usage-based enterprise pricing.



Linktree - Zero-Friction Viral Infrastructure

- **The Core Conflict:** Retaining monopolistic control over the "link-in-bio" creator ecosystem against aggressive copycat features launched by social media giants.
- **The User Retention Loop:** Massive asymmetric virality. Every time a creator puts a Linktree link in their bio, they are displaying the brand to thousands of profile visitors, creating an endless loop of organic, free user signups.
- **Product Strategy Analysis:**
 - **The Freemium Hook:** Keeps basic custom links entirely free to capture market share, while paywalling deep analytics, advanced retargeting pixels, and native checkout integrations.
 - **Aggressive Integration Ecosystem:** Deep integrations with Spotify, Shopify, and newsletter platforms turn a static link page into a central operational dashboard for independent creators.



Operator Takeaway: When your product acts as its own marketing channel, focus your engineering entirely on minimizing signup time and keeping onboarding under 60 seconds.



If I were running Linktree:

- **3 priorities:** cut signup time further; expand the analytics paywall; grow commerce integrations.
- **3 risks:** native copycats on Instagram/TikTok; low willingness to pay; churn once a creator's trend fades.
- **3 opportunities:** creator checkout and payments; audience CRM; a brand-sponsorship marketplace.



US Tech: Building the AI-Native Consumer Era



Perplexity AI vs. OpenAI - The Paradigm Shift from Chat to Synthesis

- **The Core Conflict:** Disrupting traditional search monetization by replacing standard keyword link indexes with direct, synthesized multi-source answers.
- **The User Retention Loop:** Perplexity retains power users through "Collections" and "Spaces," transforming ephemeral chat histories into organized, collaborative research assets.
- **Product Strategy Analysis:**
 - **Context Window Optimization:** Perplexity focuses its design on real-time data ingestion, routing queries to specialized web-crawlers before passing structured text to underlying LLMs, maintaining low latency.
 - **Interface Design Over Raw Compute:** While OpenAI competes on model depth and foundational intelligence benchmarks, Perplexity wins on its user-centric "Answer Engine" interface, proving that user-experience wrapping can outperform raw infrastructure.

 **Operator Takeaway:** In the AI-native generation, the winner isn't necessarily the team with the largest model - it's the team that designs the most intuitive feedback loop between human intent and machine synthesis.

 If I were Founder's Office at Perplexity:

- **3 priorities:** defend latency and answer quality; grow Spaces/Collections retention; control per-query inference cost.
- **3 risks:** Google bundling AI into core search; model-cost volatility; citation and accuracy trust gaps.
- **3 opportunities:** enterprise knowledge search; verticalized answer engines; a publisher revenue-share moat.

 Operator Toolkit

How I actually execute - grouped by function, not buzzwords.

Function	Tooling
Build	Python, SQL, FastAPI, PostgreSQL, Cursor, Claude, Copilot
Analyze	Statistics, Experimentation, Power BI, Tableau, Data Analytics
Ship	Product Scoping, MVP Development, Stakeholder Management
Operate	Process Design, Automation, Documentation, Cross-functional Collaboration

 What I'm Learning Now



Currently going deep on:

- Product Strategy
- Startup Finance
- Growth Loops
- AI Agents
- Marketplace Dynamics
- Consumer Behavior

Let's Build Together



I'm currently looking for opportunities in:

- Founder's Office
- Technical Operator
- Product & Strategy
- AI-Native Startups

If you're building something ambitious, I'd love to chat. **Let's connect.**

Shipped Products

Aa Project	📌 Role	≡ Core Metric	🔗 Live Demo
 <u>Lucid AI - Full-Stack Attention Analytics Platform</u>	Full-Stack Product Engineer	Predictive Scoring & Analytics	https://lucid-ai-self.vercel.app/
<u>VeriLoad - Solo-Shipped Dispatch Platform</u>	Solo Builder	91% Validation Accuracy	https://krrishkhanna.github.io/veriloa d/

Aa Project	📌 Role	≡ Core Metric	🔗 Live Demo
🩺 <u>SebScan - Health-Tech OCR & Risk Analytics</u>	Full-Stack Dev	87% Parse Accuracy	https://krrishkhanna.github.io/sebscan/
🔍 <u>Vachan - Multilingual Fact-Checking Platform</u>	Full-Stack Dev	+28% Precision Boost	https://vachan-1.netlify.app/landing
🧠 <u>Odin - Live Subtitles & Translation Overlay</u>	Co-Founder	100+ Languages · Freemium SaaS	https://odinsub.com/