

Krrish Khanna

Applied AI Engineer | Machine Learning | AI Systems

Chennai, India | +91 97739 34134 | krrishkhanna08@gmail.com | linkedin.com/in/krrishkhanna | github.com/krrishkhanna

SUMMARY

Applied AI Engineer building intelligent systems at the intersection of machine learning, Generative AI, and software engineering. Experienced across the full AI lifecycle—from data pipelines and deep learning models to enterprise RAG applications, LLM-powered systems, evaluation, and scalable deployment—delivering AI solutions that solve real-world business problems.

EXPERIENCE

AI Engineering Intern

Feb 2026 – Present

Git Software Technologies

Remote

- Architected production Retrieval-Augmented Generation (RAG) systems integrating document ingestion, semantic chunking, embedding generation, vector search, hybrid retrieval, and prompt orchestration for enterprise knowledge automation.
- Developed AI applications spanning computer vision, conversational AI, and Voice AI using cloud-native APIs, modular inference pipelines, and scalable deployment architectures.
- Optimized retrieval quality, latency, and response reliability through prompt refinement, retrieval tuning, and enterprise edge-case handling for CRM and customer support workflows.
- Engineered distributed AI worker infrastructure integrating PersonaPlex conversational AI and Moshi TTS to power enterprise Voice-as-a-Service deployments.

Research and Development Intern

Mar 2025 – Present

Samsung Research Institute

Bengaluru, Karnataka

- Optimized tensorized LLM inference pipelines using PyTorch, Tensor Train/Tensor Ring decomposition, and benchmarking on NVIDIA A100 infrastructure, achieving up to 1.5× inference speedup and 40% memory reduction.
- Evaluated latency-memory-quality trade-offs across compressed LLM architectures through systematic experimentation, enabling informed deployment decisions.
- Built reproducible benchmarking workflows for comparing model compression strategies under production deployment constraints.

Technical Trainee Intern – Government & Public Services

June 2025 – Aug 2025

KPMG India

Gurugram, Haryana

- Built NLP-based document intelligence pipelines combining semantic similarity, structured parsing, ontology mapping, and rule-based validation for large-scale aerospace and defence workflows.
- Automated document classification and procurement intelligence across 50K+ records per batch using Python, Sentence Transformers, SQL, and GenAI-assisted validation.
- Designed modular AI-enabled validation pipelines reducing manual review cycles and improving enterprise data quality for regulatory decision-making.

SYSTEM ARCHITECTURES & PROJECTS

Vachan – Multilingual Fact-Checking Platform | [Live Demo](#) | NLP, mBERT

Feb 2025 – Mar 2025

- Designed and deployed an end-to-end NLP pipeline integrating mBERT and the Google Fact Check API for multilingual misinformation detection, supporting 10+ languages in production.
- Applied evaluation-driven optimization cycles improving model precision by 28%; processed 15,000+ claims at 86% accuracy with a live analytics dashboard for monitoring deployment fidelity.

VeriLoad – Intelligent Validation System | [Live Demo](#) | AI Verification, REST APIs

Apr 2026 – May 2026

- Shipped a full-stack automated validation platform combining a FastAPI inference backend with a React dashboard to systematically execute dispatch logic and anomaly detection.
- Achieved 91% reconciliation accuracy across 120+ rigorous test scenarios, significantly reducing manual validation overhead through rule-based and AI-assisted processing.

SebScan – Computer Vision & OCR Analytics | [Live Demo](#) | OpenCV, Text Extraction

Apr 2026 – May 2026

- Engineered a CV and OCR pipeline achieving 87% parse accuracy across 75+ complex product labels, backed by a structured PostgreSQL schema for risk classification mapping.
- Built scalable extraction endpoints that transform raw pixel data into structured, auditable analytical outputs for downstream healthcare intelligence dashboards.

TECHNICAL SKILLS

Machine Learning: Supervised/Unsupervised Learning, Deep Learning, Model Evaluation, Feature Engineering, Statistical Learning
Generative AI: Retrieval-Augmented Generation (RAG), LLM Applications, Prompt Engineering, Agentic Workflows, Embeddings, Vector Search
AI Systems: PyTorch, FastAPI, Docker, AWS, REST APIs, Model Serving, Cloud Deployment
NLP & Computer Vision: Transformers, mBERT, Sentence Transformers, OpenCV, OCR, Semantic Similarity
Data Science: Python, SQL, Pandas, NumPy, PostgreSQL, Tableau, ETL

CERTIFICATIONS & LEADERSHIP

AWS Certified Developer – Associate | Oracle AI Vector Search Professional Certification
Machine Learning – I (Columbia University) | AWS Academy – Machine Learning Foundations
McKinsey Forward Program (Strategic Problem Solving) | **Aspire Leaders Program** (Harvard Faculty-led)